Towards A Holistic View of Bias in Machine Learning: Bridging Algorithmic Fairness and Imbalanced Learning

Damien A. Dablain · Bartosz Krawczyk · Nitesh V. Chawla

Abstract Machine learning (ML) is playing an increasingly important role in rendering decisions that affect a broad range of groups in society. ML models inform decisions in criminal justice, the extension of credit in banking, and the hiring practices of corporations. This posits the requirement of model fairness, which holds that automated decisions should be equitable with respect to protected features (e.g., gender, race, or age) that are often under-represented in the data. We postulate that this problem of under-representation has a corollary to the problem of imbalanced data learning. This class imbalance is often reflected in both classes and protected features. For example, one class (those receiving credit) may be over-represented with respect to another class (those not receiving credit) and a particular group (females) may be under-represented with respect to another group (males). A key element in achieving algorithmic fairness with respect to protected groups is the simultaneous reduction of class and protected group imbalance in the underlying training data, which facilitates increases in both model accuracy and fairness. We discuss the importance of bridging imbalanced learning and group fairness by showing how key concepts in these fields overlap and complement each other; and propose a novel oversampling algorithm, Fair Oversampling, that addresses both skewed class distributions and protected features. Our method: (i) can be used as an efficient pre-processing algorithm for standard ML algorithms to jointly address imbalance and group equity; and (ii) can be combined with fairness-aware learning algorithms to improve their robustness to varying levels of class imbalance. Additionally, we take a step toward bridging the gap between fairness and imbalanced learning with a new metric, Fair Utility, that combines balanced accuracy with fairness. Our source code and data are publicly available at https://github.com/dd1github/Fair-Over-Sampling.

Keywords Machine learning · Imbalanced data · Fairness

Damien A. Dablain and Nitesh V. Chawla Department of Computer Science and Engineering and the Lucy Family Institute for Data and Society

University of Notre Dame Notre Dame, IN 46556 USA

Bartosz Krawczyk Department of Computer Science Virginia Commonwealth University Richmond, VA 23284 USA

1 Introduction

Automated decision-making models are progressively being used in situations that affect humans in a broad range of areas, such as credit risk analysis (Khandani et al, 2010), criminal recidivism prediction (Berk et al, 2021; Chouldechova et al, 2018), hiring (Schumann et al, 2020) and the provision of social services (Shroff, 2017). For example, a bank may decide to extend credit based on whether a machine learning (ML) model predicts that an individual may default on a loan. Conversely, a judge may determine that a defendant should not be released while awaiting trial if an artificial intelligence (AI) model suggests that the defendant has a high risk of recommitting a crime. The growing prevalence of ML algorithms in decisions that affect humans is due, in part, to their perceived accuracy and ability to detect hidden patterns in data. Yet, in some cases, these models have been demonstrated to incorporate biases, such as in hiring decisions (Dastin, 2018), face recognition (Raji et al, 2020) and even translation (Caliskan et al, 2017), resulting in concerns about the fairness of machine learning algorithms (Barocas and Selbst, 2016). Despite these concerns, it is likely that the use of automated decision-making will only increase in the future, as AI becomes more wide-spread in society, government and business. Therefore, there is a growing awareness that ML algorithms should be both accurate and fair, which is underscored by the recently released Artificial Intelligence Act - a legal framework promulgated by the European Commission that mandates non-discrimination, among other requirements, for ML models that affect individuals (Commission, 2021).

Although it is desirable for ML models to be both fair and accurate, there is often a trade-off between these two goals (Kleinberg, 2018; Flores et al, 2016; Berk et al, 2021), such that increasing fairness comes at the cost of reduced accuracy. Accuracy and fairness are often at odds because they are influenced by imbalanced data. In many cases, the data used to train ML algorithms is imbalanced with respect to class and protected features, such that one class or group is underrepresented with respect to another. Since most ML models learn parameters based on data, data imbalance can cause a particular class or sub-group to be overweighted, such that preference is given to the over-represented class or group. Several previous works on algorithmic fairness do not explicitly take into account the relative sizes and imbalance of classes and protected groups (Dwork et al, 2012; Hardt et al, 2016; Zafar et al, 2017), even though these characteristics can affect accuracy and equitable outcomes. Hashimoto et al. (Hashimoto et al. 2018) has referred to the under-representation of a protected group in training data as representation disparity, such that minority groups contribute less to a ML model objective because they are under-represented in the training data, and hence model accuracy may be lower for the minority class.

Goal. To draw out the corollary between algorithmic fairness and learning from imbalanced data and introduce a novel fairness-aware oversampling approach that addresses fairness within the construct of learning from imbalanced data.

Summary. The algorithmic fairness domain focuses on combating bias in decision-making originating in protected features that could affect the objectiveness of the decision and lead to unfavorable results for under-represented and minority groups. At the same time, the class imbalance domain focuses on countering bias originating from skewed class distributions, as majority classes may be preferred

over minority ones during classifier training. We postulate that when handling protected features to achieve fairness, one will always face some kind of imbalance in the training data. Therefore, algorithmic fairness and class imbalance should be seen as two sides of the same coin. Using techniques from the imbalanced data domain can benefit algorithmic fairness, leading to the design of more robust and holistic methods that can counter both data and algorithmic biases. We take a step toward bridging the gap between algorithmic discrimination and imbalanced learning by discussing the key concepts and metrics that underpin both areas. Because it is often not possible for a ML algorithm to meet multiple fairness criteria (e.g., individual and group non-discrimination) at the same time (Chouldechova, 2017; Corbett-Davies et al, 2017), we focus on a single element - group fairness. We show that common approach used in imbalanced learning - data oversampling - can be used to increase model fairness and accuracy.

Main contributions. This paper offers the following insights and contributions towards fair and robust machine learning:

- Bridging the Gap Between Fairness and Imbalanced Learning: We take a step toward bridging the gap between the algorithmic fairness and imbalanced learning fields by showing that unbalanced training datasets affect a classifier's ability to produce equitable treatment of protected groups. Training sets that are proportionately balanced with respect to classes and protected groups result in fairer results when measured by three measures of group fairness average odds, equalized odds and the difference in true negative rates.
- Fair Oversampling: We propose a new data pre-processing technique, Fair Oversampling (FOS), that enhances fairness and classifier accuracy. The method is based on a popular imbalanced learning technique that is adaptable to group fairness. It is straight-forward and can be easily inserted into a fairness-aware ML model or standard classifier pipeline, improving robustness under varying levels of class imbalance.
- Fair Utility metric: We propose a new metric that combines fairness with imbalanced learning - Fair Utility - that relies on measures commonly used in both fields.

Organization. The paper is organized as follows. It first discusses algorithmic fairness and its key concepts. Then, it reviews the central elements of imbalanced learning. Next, it introduces our algorithm and proposed fairness metric. Finally, the paper discusses experimental results, and the commonalities between algorithmic fairness and imbalanced learning.

2 Algorithmic fairness

Discrimination can generally be defined as the prejudicial treatment of an individual based on membership in a legally protected group. Algorithmic fairness is concerned with ensuring that decisions made by machine learning models are equitable with respect to protected groups (Romei and Ruggieri, 2014; Žliobaite, 2017). Algorithmic fairness commonly invokes one of the following changes in order to produce equitable results: (1) modifications to the training data, (2) changes in the machine learning model, or (3) modifications to the decisions themselves (Calmon et al, 2017).

We concentrate on supervised learning in the context of binary classification. In binary classification, the goal of algorithmic fairness is to fairly select between two actions, a_0 and a_1 (e.g., approve or decline the extension of credit in banking). In our discussion, we adopt notation used by Speicher et al. (Speicher et al, 2018) to describe our algorithmic environment. Thus, a ML decision algorithm, A, can be described as a function $A: \mathbb{R} \to \{0,1\}$ that outputs a binary decision. The machine learning algorithm, A, parameterized by θ , accepts as input training data, D, minimizes a loss function $l(\theta)$, and predicts a label (i.e., 0 or 1).

More formally, A accepts as input training data, $D = \{(x_i, y_i)\}_{i=1}^n$, with n examples, features $x_i \in X$, where $y_i \in Y$ represents the prediction or label $(Y = \{0, 1\})$ for each individual i. The features, X, can either be discrete or continuous. We partition the set of features (or attributes) into two groups, sensitive or protected features, such as gender, race or age, and unprotected features, such that $x = (x_p, x_u)$. We also assume that protected features can be further partitioned into two classes, privileged and unprivileged (i.e., $x_p = (x_{pr}, x_{up})$). For purposes of this paper, we assume that the label contained in the dataset is the correct, unbiased label.

Narayanan described at least 21 mathematical definitions of fairness that have been proposed by the fairness research community (Narayanan, 2018). Two broad classes of algorithmic equity have gained prominence: group and individual fairness. Group fairness requires that the ML algorithm, A, produces parity for a given metric, M, for protected features, such that $M_{x_p}(A) = M_{x_p'} \forall x_p, x_p' \in X$. Individual fairness requires that similar individuals are treated similarly. It implies the presence of a similarity metric that is capable of determining if a pair of individuals are similar. The leading definition of individual fairness is metric fairness, as proposed by Dwork et al. (Dwork et al, 2012). It requires: $M_y(A(x_1), A(x_2)) \leq LM_x(x_1, x_2) \ \forall \ x_1, x_2 \in X$, where M_x and M_y are data and label metrics, A is a map or algorithm, and $L \geq 0$ is a Lipschitz constant. Finding such a similarity metric can be challenging. To overcome this obstacle, some works have assumed the existence of a publicly available similarity metric issued by a regulatory body that follows Rawlsian principles (Rawls, 2001; Dwork et al, 2012).

Because of the challenges in finding suitable individual fairness similarity metrics, we focus on group fairness in this paper. Corbett-Davies et al. describe three central concepts embodied by group fairness: anti-classification, classification parity and calibration (Corbett-Davies and Goel, 2018). Anti-classification requires that AI algorithms do not consider protected features when making decisions (Bonchi et al, 2017; Caliskan et al, 2017; Grgic-Hlaca et al, 2016). Thus, anti-classification provides that: $A(x) = A(x') \ \forall x$, such that $x_u = x'_u$. Classification parity (sometimes referred to as statistical parity) requires that certain measures are equal across sensitive features. Statistical parity can be expressed in a variety of ways. Under one formulation, the proportion of members in a protected group receiving a positive classification must be identical to the proportion in the population as a whole (Zemel et al, 2013). Other measures focus on the difference in positive or negative rates (instead of proportions) between sensitive groups (e.g., equal true positive rates for both male and female applicants). Classification parity has been widely used as a fairness metric in machine learning (Agarwal et al, 2018; Calders and Verwer, 2010; Edwards and Storkey, 2015). As discussed below, we use classification parity in our metrics. Demographic parity, or the proportion of positive decisions, means that

 $Pr(A(X) = 1|x_p) = Pr(A(X) = 1)$ (Feldman et al, 2015). Whereas, parity of false positives requires that $Pr(A(X) = 1|Y = 0, x_p) = Pr(A(X) = 1|Y = 0)$. We also incorporate demographic parity into our metrics, although we focus on differences in true positive, false positive, and true negative rates, instead of their relative proportions.

Group fairness concepts are firmly entrenched in U.S. society and law. For example, the Fourteenth Amendment to the U.S. Constitution contains an equal protection law that prevents government workers from acting with "discriminatory purpose" against protected groups. Laws issued by the US Supreme Court also prevent discriminatory actions against protected groups, as formalized by *Griggs v. Duke Power Co.*. The ubiquity of group fairness concepts in society is another reason why we focus on them in our analysis.

Fairness-aware algorithms. In order to achieve group fairness in machine learning, a variety of techniques have been employed, which can be broadly separated into pre-processing, in-processing and post-processing methods. Pre-processing techniques involve manipulating the training data before it is consumed by a classification algorithm, in-processing incorporates fairness into a ML algorithm loss function, and post-processing aims to adjust the decisions of a classifier to be fair. We briefly survey below the key pre-processing and in-processing techniques that are relevant to our approach.

Pre-processing techniques. Kamiran and Calders propose a pre-processing method, Reweighing, that creates weights for the training instances to ensure fairness (Kamiran et al, 2013). They effectively divide the training set into four groups: (1) privileged group, majority class; (2) unprivileged group, majority class; (3) privileged group, minority class; and (4) unprivileged group, minority class. They then develop separate weights for each of the four groups and apply the weights to each instance. Feldman et al. propose a pre-processing method, Disparate Impact Remover, that modifies features to enhance group fairness while preserving rank-ordering within protected groups (Feldman et al, 2015).

In-processing techniques. Agarwal et al. propose an in-processing method, Exponentiated Gradient Reduction (EGR), that reduces fair classification to a sequence of cost-sensitive classification problems, whose solution yields a randomized classifier with the lowest error subject to the desired constraints (Agarwal et al, 2018). Zhang et al. propose Adversarial Debiasing to mitigate unwanted biases (Zhang et al, 2018). Their method uses adversarial training to prevent an adversary from determining a protected attribute from a ML model's predictions. This approach results in a fair classifier because the predictions are not allowed to include group discrimination information. Yurochkin, Bower and Sun propose a method, Sensitive Subspace Robustness (SSR), based on adversarial training (Yurochkin et al, 2019). SSR induces individual fairness based on sensitive perturbations of inputs. It casts fairness in the form of robustness to sensitive perturbations of the training data. SSR uses a fair Wasserstein distance metric to require the output of a ML model to be similar to the training label.

3 Imbalanced learning

Imbalanced learning is concerned with disproportions among classes. In binary classification, the number of instances of one class (the majority) outnumber the

other (minority). The skewed distribution of examples in favor of the majority class can cause classifiers to be biased toward the majority because the algorithm's parameters are more heavily weighted toward more frequently occurring examples. Classifiers can achieve high accuracy by merely selecting the majority class. However, the minority class is often the more important one from the data inference perspective because it may carry more relevant information. For example, in anomaly detection, there may be few examples of malicious software, which is the more interesting class. In addition to detecting malicious software in computer security (Cieslak et al, 2006), imbalanced learning has been applied to fraud detection (Wei et al, 2013), medical treatment (i.e., identifying rare diseases) (Rao et al, 2006), image recognition (Kubat et al, 1998), atypical behavior in social networks and infrastructure monitoring systems (Krawczyk, 2016). Training sets where one class is under-represented compared to the majority class frequently results in poor accuracy performance by the classifier (Seiffert et al, 2007).

Taxonomy. There are three broad approaches within imbalanced learning: datalevel methods that modify the training data to balance class distributions, algorithm level methods that ameliorate bias in classifiers towards the majority class, and ensemble methods that are a combination of the first two with classifier committees.

Data-level approaches. This group of methods focus on modifying the training set by balancing the number of minority and majority class examples. Oversampling generates new minority class examples, while under-sampling removes instances from the majority class. Under-sampling can result in removal of important data from the training set and therefore is often not preferred. Simple random oversampling (ROS) merely duplicates instances of the minority class to impose parity. SMOTE, or the Synthetic Minority Oversampling Technique, (Chawla et al, 2002), is a popular oversampling method used in the imbalanced learning community. It randomly selects a nearest neighbor of a minority instance and linearly generates synthetic examples based on the original instance and a nearest neighbor. SMOTE has been adapted to enhance the importance of class borderline instances (Han et al. 2005), define safe regions that do not sample from noisy or overlapping instances (Bunkhumpornpat et al, 2009) and has been applied in the deep learning (Dablain et al, 2022) and big data (Sleeman and Krawczyk, 2021) contexts. Alternative approaches to SMOTE have been proposed recently that do not rely on k-nearest neighbors, instead using alternative measures such as class potential (Krawczyk et al, 2020), Mahalanobis distance (Sharma et al, 2018), or manifold approximation (Bej et al, 2021).

Algorithm-level approaches. This group of methods modify the training procedure of a classifier to make it skew-insensitive, or incorporate alternative cost functions. Cost sensitive learning, which is a form of importance sampling (Kahn and Marshall, 1953), magnifies the importance of minority examples by increasing the penalty associated with the instances. Recent examples of cost-sensitive methods that have been used in imbalanced learning include the focal loss (Lin et al, 2017), the class-balanced margin loss (Cui et al, 2019), the distribution aware margin loss (Cao et al, 2019) and the asymmetric loss (Ridnik et al, 2021).

Ensemble approaches. Combining multiple classifiers is considered as one of the most effective approaches in modern machine learning (Woźniak et al, 2014). Ensembles find their natural application in learning from imbalanced data, as they

leverage the predictive power of multiple learners. By combining base classifiers with data or algorithm-level solutions, they achieve locally specialized robustness and maintain diversity among ensemble members. Most popular solutions combine resampling with Bagging (Lango and Stefanowski, 2018) or Boosting (Zhang et al, 2019b), use mutually complimentary cost-sensitive learners (Tao et al, 2019), or rely on dynamic selection mechanisms to tackle locally difficult decision regions (Zyblewski et al, 2021).

4 Why we need to bridge algorithmic fairness and imbalanced learning?

Different views on bias. The two previous sections provided general background and reviewed recent advancements in algorithmic fairness and imbalanced learning. This allows us to see the strong parallel between them, as they both deal with the problem of countering bias in data and algorithms, however from different perspectives:

- Bias according to algorithmic fairness. Here, bias is seen as a lack of fairness and transparency, originating from social background and the nature of the data itself. Fairness focuses on bias based on using sensitive or protected information (e.g., race or gender) to make a decision, thus putting traditionally under-represented groups at a significant disadvantage. Fairness-aware algorithms also focus on using safe information for training classifiers and debiasing them with respect to protected features.
- Bias according to imbalanced learning. Class imbalance focuses on bias originating in disproportion among classes, as most machine learning algorithms will become biased towards classes with a higher number of training instances. This puts smaller, yet often more important, classes at a disadvantage. Imbalance-aware algorithms focus on either balancing class distributions or removing the bias towards majority classes from the training process.

Two sides of the same coin. We can see that algorithmic fairness views the source of bias as the information being used to make a decision, while class imbalanced learning views bias as arising from a disproportion among instances used to train a classifier. In this paper, we argue that these are two sides of the same coin. While bias in decision making may stem from using sensitive features, they are usually followed by imbalanced distributions, as traditionally under-represented groups are akin to minority classes. Therefore, considering only one type of bias leads to an oversimplified view of the problem. To achieve a truly fair and robust machine learning algorithms, we need to develop a holistic view where both types of bias are tackled during classifier training.

5 Fair Oversampling

Our algorithm, Fair Oversampling (FOS), is designed to improve fairness and increase classifier accuracy. The basic intuition behind FOS is that bias in machine learning models is caused, in part, by under-represented classes and features in training sets. When training a machine learning model to accurately and fairly

discern classes and protected features, it is often necessary to equalize the *number* of training examples between classes and protected groups to ensure that algorithmic models based on parametric learning are able to balance gradients and model weights between specific classes and features. If a classifier observes very few instances of a minority class or certain minority features, it's parameters may be biased toward recognizing the dominant class or group. FOS recognizes that numerical class and feature imbalance may exist in the data used to train machine learning models and aims to redress this numerical data imbalance - thus enhancing both model predictive accuracy and fairness.

As a pre-processing algorithm, it modifies a training dataset D so that it can be input to the machine learning model. The modified data can either be used with a standard classifier, such as Support Vector Machines (SVM), or with an algorithm that is specifically designed to improve fairness, such as SSR or Adversarial Debiasing. FOS acts on two types of independent variables (X) in the training data, protected features x_p , which are binary, and unprotected features x_u , which are expressed as real numbers. The pseudocode for FOS is displayed in Algorithm 1.

FOS first determines the minority and majority classes $(Y = \{0,1\})$ or $Y = \{min, maj\}$). It then subdivides the protected features x_p into two categories - privileged and unprivileged $(x_p = (x_{pr}, x_{up}))$. For example, in a dataset D related to the extension of credit, the majority class (D_{maj}) could be people that receive credit, and the minority class (D_{min}) could be those that do not receive credit. The protected feature x_p could be gender, where males are considered privileged x_{pr} and females are unprivileged x_{up} . This categorization results in four subgroups: privileged majority (D_{prmaj}) , unprivileged majority (D_{upmaj}) , privileged minority (D_{prmin}) and unprivileged minority (D_{upmin}) .

The objective of FOS is to restore balance between the classes and protected features through random oversampling and nearest neighbor metrics, using the mechanics of the SMOTE algorithm, but with modifications. FOS numerically balances the classes and protected features, such that the number of examples (N) in the majority class (N_{maj}) equals the number of examples in the minority class (N_{min}) , or $N_{maj}=N_{min}$. It determines the protected group x_p in the dataset D that requires the least number of samples to obtain equivalency (denoted as D_1), and selects the K nearest neighbors of a random sample of D_1 (e.g., the unprivileged, minority group D_{upmin}). In our implementation, five nearest neighbors were selected; however, K selection can depend on the specific dataset. The number of random samples selected from this group equals the number of samples required to make it equal in number to the same group in the majority class (e.g., $N_{upmin}=N_{upmaj}$). For D_1 , the samples are drawn from a single protected sub-group (e.g., the samples are exclusively drawn from D_{upmin}).

Next, the same oversampling procedure is repeated for the protected group x_p with the larger number of samples that are required to obtain numerical equivalency (e.g., D_{prmin}), which is denoted as D_2 , except that instead of drawing the nearest neighbors exclusively from the D_2 pool, they are drawn from the entire minority class. For example, the samples are drawn from D_{min} instead of D_{prmin} . Empirically, we found that this approach reduced bias because it blurs the difference between privileged minority D_{prmin} and unprivileged minority D_{upmin} group members, since it draws a nearest neighbor from the entire minority class D_{min} , which consists of both privileged and unprivileged members. This approach

Algorithm 1: Fair Oversampling Pseudocode

```
Notation:
D - training dataset;
L - training labels;
D_{prmaj} - privileged group, majority class; D_{upmaj} - unprivileged group, majority class;
D_{prmin} - privileged group, minority class;
\hat{D_{upmin}} - unprivileged group, minority class;
egin{aligned} N_{prmaj} & - & \text{number of } D_{prmaj}; \ N_{upmaj} & - & \text{number of } D_{upmaj}; \ N_{prmin} & - & \text{number of } D_{prmin}; \end{aligned}
N_{upmin} - number of D_{upmin};
 S_{pr} = number of D_{prmin} to oversample;
\hat{S}_{up} = \text{number of } \hat{D}_{upmin} \text{ to oversample};
Calculate number of instances to oversample:
S_{pr} = N_{prmaj} - N_{prmin};

S_{up} = N_{upmaj} - N_{upmin};

if S_{up} is less than S_{pr} then

N_{samp1} = S_{up};
       N_{samp2} = S_{pr};

D_1 = D_{upmin};
      D_2 = D_{prmin};
_{
m else}
       N_{samp1} = S_{pr};
      N_{samp2} = S_{up};

D_1 = D_{prmin};
      D_2 = D_{upmin};
end
Base_1 \leftarrow randomly \ select \ N_{samp1} \ from \ D_1;
for B in Base_1 do
       KNN \leftarrow select \ K \ nearest \ neighbors \ of \ B \ from \ D_1;
       N \leftarrow randomly \ select \ 1 \ neighbor \ from \ KNN;
       R \leftarrow randomly \ select \ a \ number \ between \ 0 \ and \ 1; S = B \ + \ R \ X \ (B \ - \ N);
       Samples \leftarrow append S;
       Labels \leftarrow append\ label\ of\ D_1;
end
X \leftarrow D + Samples;
Y \leftarrow L + Labels;
Base_2 \leftarrow randomly \ select \ N_{samp2} \ from \ D_2; for B \ in \ Base_2 \ \mathbf{do}
       KNN \leftarrow \text{select } K \text{ nearest neighbors of } B \text{ from } D_2;
       N \leftarrow randomly \ select \ 1 \ neighbor \ from \ KNN;
       R \leftarrow randomly \ select \ a \ number \ between \ 0 \ and \ 1; S = B \ + \ R \ X \ (B \ - \ N);
       Samples \leftarrow append \ S;
      Labels \leftarrow append\ label\ of\ D_1;
\mathbf{end}
X \leftarrow X + Samples;
Y \leftarrow Y + Labels
```

provides an opportunity to effectively increase the representation of privileged, minority D_{prmin} members if the features of a privileged minority D_{prmin} member are nearest in proximity to an unprivileged, minority member D_{upmin} . Our sampling procedure is different than the SMOTE approach, which selects nearest neighbors solely from the combined minority class, without regard to protected groups.

FOS balances the *number* of class and protected group instances within a training dataset. It does not balance protected attribute *ratios*. We postulate that FOS produces fair results, even without balanced protected group ratios, because it acts on the weights of parametric machine learning models. By allowing the machine learning models to observe an equal *number* of class and protected group

instances, it facilitates balanced parametric learning (e.g., equal weighted average number of instances), even though the natural ratio of feature instances may vary.

6 Experiments

The following experiments were designed to answer the following research questions:

- RQ1: Does FOS improve both algorithmic fairness and robustness to class imbalance for popular standard classifiers?
- RQ2: Can FOS further improve the performance of fairness-aware classifiers?
- RQ3: What is the FOS trade-off between fairness and skew-insensitive metrics when handling varying imbalance ratio levels?
- RQ4: How does FOS influences feature importance used by underlying the classifiers and does it lead to more fairness-aware feature selection?

6.1 Datasets

Three popular datasets were selected for testing that were used by the fairness research community (Quy et al, 2021): German Credit (Hofmann, 1994), Adult Census Income (Kohavi et al, 1996), and Compas Two-Year Recidivism (Angwin, 2016). The key statistics of each dataset are summarized in Table 1. All three datasets involve binary classification. The German Credit dataset contains data that allows a classifier to predict whether an individual should have a positive or negative credit rating. The Adult Census Income dataset predicts whether an individual earns more or less than \$50K. The Compas dataset can be used to predict whether a defendent will commit a crime within a two year period.

As we can see from Table 2, all of the datasets exhibit both class and protected attribute (gender) imbalance. Compas shows the least amount of class imbalance, with a ratio of 1.22:1, while the German Credit and Adult Census datasets have class imbalance ratios ranging from approximately 2:1 to 3:1. All three datasets show greater protected attribute imbalance than class imbalance, with the ratios ranging from approximately 2:1 to 4:1. In the minority class, the maximum protected attribute imbalance ratios are even higher, ranging from 1.75:1 in German Credit to 5.61:1 in Adult Census.

Table 1: Description of the Datasets

Dataset	Instances	Features	Protected Feature	Classes
German Credit	1,000	20	Gender	Good credit; bad credit
Adult Census Income	48,842	14	Gender	Income > or < 50K
Compas	7,214	28	Gender	Recidivism; No recidivism

Table 2: Class and protected feature imbalance ratios for each dataset

Class / Feature	German Credit	Adult Census	Compas				
Classes							
Majority	700	37,155	3,963				
Minority	300	11,687	3,251				
Ratio Protected Features	2.33	3.18	1.22				
Privileged	690	32,650	5,819				
Unprivileged	310	16,192	1,395				
Ratio	2.23	2.02	4.17				
Combined Class and Protected Features							
Privileged, majority	499	22,732	3,066				
Unprivileged, majority	201	14,423	897				
Ratio	2.48	1.58	3.42				
Privileged, minority	191	9,918	2,753				
Unprivileged, minority	109	1,769	498				
Ratio	1.75	5.61	5.53				

6.2 Experimental design

Experiment 1: Oversampling for standard classifiers. First, modified training data produced by FOS was used as input to two standard machine learning classifiers: SVM and Logistic Regression (LG). The performance of the models was assessed based on metrics, which are discussed below. The performance of our algorithm was compared against four benchmarks for each standard classifier: (1) a baseline (no modifications to the training dataset); (2) a popular imbalanced learning oversampling method - SMOTE (Chawla et al, 2002); and two pre-processing algorithms that are specifically designed to improve fairness - (3) Reweighing (Calmon et al, 2017) and (4) Disparate Impact Remover (Feldman et al, 2015). The purpose of this experiment was to determine how FOS compared to other data pre-processing algorithms that are used in both the imbalanced learning and fairness research communities.

Experiment 2: Oversampling for fairness-aware classifiers. Second, we assessed whether FOS could be used in conjunction with in-processing fairness-aware algorithms. For this experiment, the SSR (Yurochkin et al, 2019), Adversarial Debiasing (Zhang et al, 2018) and EGR models (Agarwal et al, 2018) were chosen. To determine whether FOS could improve accuracy and fairness for these algorithms, the models were alternatively run with data that was not pre-processed and data that was modified by FOS.

Experiment 3: Robustness to increasing imbalance ratios. Third, we assessed how the performance of a standard ML classifier was affected by increasing levels of class and protected group imbalance. For this test, SVM was used as the ML algorithm with varying degrees of imbalance. Instances were randomly removed from classes and protected groups to achieve the intended imbalance levels. The selected imbalance levels were: $I \in \{1, 1.2, 1.4, 1.6, 1.8, 2\}$ for German Credit; $I \in \{1, 2, 2.5, 3, 3.5, 4\}$ for Compas; and $I \in \{1, 2, 4, 6, 8, 10\}$ for Adult Census,

where I represents the denominator of the fraction that reduced the number of original protected group members. The reason for different levels of imbalance by dataset is that classifiers trained with the Reweighing and Disparate Impact Remover algorithms produced unstable results for datasets with a relatively small number of examples (i.e., German Credit and Compas), such that the classifiers predicted all labels to reside in a single class, thus causing True Negatives to be zero, yielding "Nan" metrics. In contrast, both SMOTE and FOS were able to work at the I=10 level for all datasets. Therefore, the imbalance ratio scaling was adjusted so that all pre-processing algorithms could be assessed for all datasets.

Experiment 4: Impact of oversampling on feature importance. Fourth, we considered whether FOS caused a standard classifier to change the selection of the features that it used to formulate its decision boundary. For this purpose, the SVM ML algorithm was selected and feature importances of a classifier were compared using baseline training data and data modified by FOS to determine whether modifying the training data changed model feature selection.

Setup. All experiments were performed using five fold cross-validation. The reported results are averaged over the respective held out validation sets. See Tables 3, 4, 5.

6.3 Metrics

For purposes of the experiments, group fairness metrics were selected that could be expressed as elements of a binary classification confusion matrix consisting of True Positive Rate (TPR), False Positive Rate (FPR), True Negative Rate (TNR), and False Negative Rate (FNR). Metrics were chosen that are widely used in the fairness and imbalanced learning communities: Balanced Accuracy, Average Odds Difference (AOD), Absolute Average Odds Difference (AAO), Equal Opportunity Difference (EOD), and True Negative Rate Difference (TNRD) (Hardt et al, 2016; Biswas and Rajan, 2020; Chakraborty et al, 2021). AOD is the average difference in the False Positive Rate plus the True Positive Rate for privileged and unprivileged groups (Bellamy et al, 2018). It can be expressed as: $\frac{1}{2}((TPR_p - TPR_{up}) + (FPR_p - FPR_{up}))$, where TPR_p is the TPR of privileged instances, TPR_{up} is the TPR of the unprivileged instances, FPR_p is the FPR of privileged instances, and FPR_{up} is the FPR of the unprivileged instances. AAO is the same as AOD, except that TPR and FPR are absolute value calculations. EOD is the difference between the True Positive Rate of privileged and unprivileged groups (Bellamy et al, 2018), and can be expressed as: $(TPR_p - TPR_{up}) + (FPR_p - FPR_{up})$. TNRD is $(TNR_p - TNR_{up})$, where TNR_p is the TNR of privileged instances, and TNR_{up} is the TNR of the unprivileged instances.

6.3.1 Proposed metric.

In addition to the metrics discussed above, we propose a new metric, called **Fair Utility**. In developing this metric, we are inspired by Corbett-Davies et al. (Corbett-Davies et al. (2017). They characterize algorithmic fairness in terms of constrained optimization in the context of the COMPAS algorithm for determining whether defendants in Broward County, FL, who were awaiting trial, were

too dangerous to be released. In their formulation, the objective of algorithmic fairness is to both maximize public safety and reduce racial disparities. We also view algorithmic fairness as a multi-objective optimization problem, where the goal is to maximize the accuracy of a classifier and reduce group inequality. We approach the optimization problem with a data pre-processing technique designed to balance class accuracy prediction with protected group equity. We are also inspired by Halevy, Norvig and Pereira, who postulated that, in machine learning, a large quantity of data is more important than a strong algorithm (Halevy et al, 2009). Our metric named Fair Utility can be expressed as balanced accuracy multiplied by the average of TPRD plus TNRD. More explicitly, it is: $BA X \frac{1}{2} X ((1 - |TPRD|) + (1 - |FPRD|))$, where BA is balanced accuracy, TPRD is $(TPR_p - TPR_{up})$, and FPRD is $(FPR_p - FPR_{up})$. Utility involves maximizing the benefit of taking an action, compared with its costs. Here, we treat accuracy as equivalent with utility, which assumes that the class label assigned by the dataset is correct and does not contain inherent bias. The objective of Fair Utility is to combine accuracy and fairness into a single metric by incorporating balanced accuracy (which reflects the impact of class imbalance) with two fairness metrics (true positive and true negative rates) that track whether a classifier consistently accepts or rejects protected group members.

6.4 Experiment 1: Oversampling for standard classifiers

FOS displays strong performance with respect to the standard classifiers, with clear improvements in fairness, as measured by AOD, AAO, TNRD, and EOD; although it shows better results with SVM as compared to Logistic Regression. For the SVM classifier, it consistently outperforms the other algorithms in terms of average odds, absolute average odds, TNRD, and Fair Utility. See Table 3. It came in a close second to SMOTE in terms of Balanced Accuracy and clearly outperformed SMOTE in terms of the fairness metrics. Although SMOTE displays strong balanced accuracy, it often does not produce fair results with respect to protected groups. This is likely because it balances the class distribution, which improves the class false positive rate; while it is not designed to improve the false positive rate with respect to specific instance features. In terms of equal odds, FOS demonstrates significant reductions in unfairness compared to the baseline, with a first place finish for Compas and second place finishes for German Credit and Adult Census.

For Logistic Regression, FOS consistently produced the top Fair Utility results, with first place finishes in terms of average odds and absolute average odds, and first place misses by less than .0057 points. It also showed significant reductions in equal odds and TNRD, when compared to baselines, with first or second place results. See Table 4.

For purposes of this experiment, FOS consistently demonstrates that it improves both accuracy and fairness over baselines. It also outperforms other fairness pre-processing algorithms on a number of measures. Thus, this experiment shows that an oversampling technique that is adopted from imbalanced learning can achieve significant improvements in group fairness measures. This also shows the close relationship between class and protected group imbalance and fairness - by jointly improving class and protected group imbalance ratios, we can affect

Table 3: SVM Classifier Results

Method	Bal Acc	Avg Odds	Absol Avg Odds	TNRD	Equal Odds	Fair Utility
German Dataset						
Base	0.6486	0.1054	0.1054	0.1358	0.0750	0.5795
SMOTE	0.7055	0.0500	0.0531	0.0812	0.0250	0.6677
Reweigh	0.6883	0.0411	0.0492	0.0506	0.0478	0.6546
Disparate	0.6152	0.0745	0.0788	0.1071	0.0504	0.5668
FOS	0.7003	0.0174	0.0313	0.0166	0.0460	0.6781
Adult Dataset						
Base	0.7052	0.1886	0.1886	0.0666	0.3106	0.5722
SMOTE	0.8044	0.3094	0.3094	0.2934	0.3255	0.5527
Reweigh	0.7667	0.0626	0.0639	0.1174	0.0103	0.7177
Disparate	0.7259	0.4378	0.4378	0.2687	0.6070	0.4035
FOS	0.7935	0.0208	0.0422	0.0247	0.0598	0.7600
Compas Dataset						
Base	0.6641	0.2022	0.2022	0.1541	0.2503	0.5296
SMOTE	0.6700	0.2176	0.2176	0.1835	0.2517	0.5242
Reweigh	0.6447	0.0319	0.0353	0.0462	0.0244	0.6220
Disparate	0.6613	0.0486	0.0500	0.0593	0.0407	0.6280
FOS	0.6680	0.0125	0.0246	0.0267	0.0224	0.6512

Table 4: Logistic Regression Classifier Results

Method	Bal Acc	Avg Odds	Absol Avg Odds	TNRD	Equal Odds	Fair Utility
German Dataset						
Base	0.6469	0.0886	0.1084	0.0730	0.1438	0.5756
SMOTE	0.7158	0.0659	0.0867	0.0688	0.1046	0.6536
Reweigh	0.6099	0.0679	0.0679	0.0387	0.0970	0.5683
Disparate	0.6385	0.0648	0.0793	0.0417	0.1169	0.5875
FOS	0.7144	0.0452	0.0563	0.0414	0.0712	0.6746
Adult Dataset						
Base	0.6765	0.1960	0.1960	0.0726	0.3193	0.5439
SMOTE	0.7595	0.3238	0.3238	0.3776	0.2700	0.5136
Reweigh	0.6941	0.0080	0.0113	0.0056	0.0170	0.6862
Disparate	0.6532	0.0359	0.0359	0.0156	0.0561	0.6298
FOS	0.7352	0.0137	0.0158	0.0162	0.0153	0.7236
Compas Dataset						
Base	0.6649	0.2129	0.2129	0.2704	0.1553	0.5230
SMOTE	0.6733	0.2271	0.2271	0.2578	0.1963	0.5200
Reweigh	0.6669	0.0170	0.0285	0.0327	0.0243	0.6476
Disparate	0.6675	0.0844	0.0844	0.1027	0.0661	0.6111
FOS	0.6674	0.0139	0.0277	0.0244	0.0309	0.6487

a substantial improvement in group fairness measures. These results also indicate that it is possible to increase both balanced accuracy and fairness simultaneously ($\mathbf{RQ1}$ answered).

6.5 Experiment 2: Oversampling for fairness-aware classifiers

In addition, FOS regularly improves the accuracy and fairness metrics of fairness-aware algorithms (see Table 5). On the Compas dataset, it uniformly improves results, and in cases where it does not (e.g., equal odds), it misses first place by .0013. In the case of Adult Census, it generally displays better results, and in cases where there is a slight degradation in fairness, there is a substantial improvement

Table 5: Fairness Classifiers With & Without FOS

Method	Bal Acc	Avg Odds	Absol Avg Odds	TNRD	Equal Odds	Fair Utility
German Dataset Exp. Grad. Exp. Grad. (+FOS)	0.6562 0.7160	0.0716 0.0306	0.0857 0.0437	0.0557 0.0271	0.1156 0.0602	0.5993 0.6845
Adver. Debias. Adver. Debias. (+FOS)	0.5739 0.6680	0.6006 0.3988	0.6006 0.3988	0.6143 0.3989	0.5868 0.3986	0.2613 0.4272
SSR SSR (+FOS)	0.6999 0.7026	0.0032 0.0677	0.0545 0.0720	0.0577 0.0459	0.0513 0.0980	0.6617 0.6528
Adult Dataset Exp. Grad. Exp. Grad. (+FOS)	0.6676 0.7387	0.0089 0.0107	0.0089 0.0133	0.0114 0.0154	0.0063 0.0112	0.6617 0.7288
Adver. Debias. Adver. Debias. (+FOS)	0.7466 0.8132	0.0528 0.0430	0.0594 0.0671	0.0662 0.0241	0.0526 0.1101	0.7022 0.7586
SSR SSR (+FOS)	0.6976 0.7132	0.1047 0.0293	0.1047 0.0325	0.1327 0.0500	0.0768 0.0151	0.6246 0.6901
Compas Dataset Exp. Grad. Exp. Grad. (+FOS)	0.6638 0.6697	0.0167 0.0147	0.0327 0.0300	0.0379 0.0308	0.0276 0.0293	0.6419 0.6492
Adver. Debias. Adver. Debias. (+FOS)	0.6718 0.6775	0.2066 0.1267	0.2006 0.1267	0.2291 0.1345	0.1720 0.1190	0.5369 0.5922
SSR SSR (+FOS)	0.6251 0.6521	0.0649 0.0142	0.0649 0.0184	0.0750 0.0161	0.0549 0.0207	0.5848 0.6440

in accuracy. Thus, FOS can also supplement existing fairness-aware algorithms as a data pre-processing step ($\mathbf{RQ2}$ answered).

6.6 Experiment 3: Robustness to increasing imbalance ratios

FOS performs at the top of the benchmark group in terms of Balanced Accuracy and Fair Utility under increasing levels of imbalance on all three datasets. See Figure 1. However, at first glance, it does not outperform in terms of discrimination mitigation at higher levels of imbalance. Upon closer inspection, we believe that the reason why the baseline and other algorithms appear to be more stable at higher imbalance ratios is because their predictions focus on the true positives at the expense of true negatives. This can be seen in the Adult Census and Compas datasets, which have higher levels of imbalance. In those cases, as depicted in Figure 2, the precision ratios increased and the recall ratios decreased for most algorithms, except for FOS and SMOTE (RQ3 answered). As discussed in the Experiments section, it should be remembered that other pre-processing techniques initially failed at imbalance levels greater than 2 and 4 on the German Credit and Compas datasets, respectively.

6.7 Experiment 4: Impact of oversampling on feature importance

For the German Credit and Adult datasets, FOS induces some changes in the classifier's feature selection and feature importance. For Compas, FOS does not trigger much of a change in feature selection, likely because it has fewer features than the other two datasets. See Figure 4. For German Credit, FOS increases the importance of the checking account feature (from .06 to .10) and moves the debtors

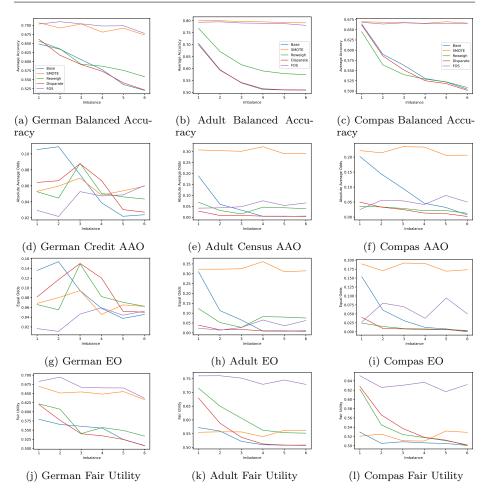


Fig. 1: Impact of Varying Imbalance Levels. This figure illustrates the impact of varying protected group, minority class imbalance ratios on key performance metrics using a SVM classifier. It compares the performance of the baseline (no training set modifications), 3 benchmark algorithms (SMOTE, Reweighing, and Disparate Impact, and FOS. FOS shows high resilience to increasing imbalance levels with respect to Balanced Accuracy and Fair Utility; while exhibiting less stability with respect to AAO and EO. We hypothesize that the reason for this seeming instability is that the baseline and other benchmarks focus (or favor) majority class predictions over minority class predictions; whereas FOS maintains an even balance between majority and minority class predictions. We can see this in the next figure.

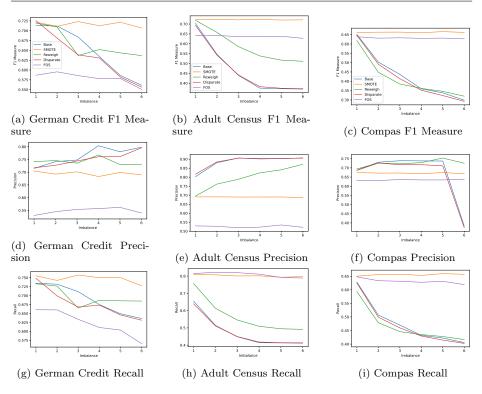


Fig. 2: Varying Imbalance Levels on F1, Precision & Recall. This figure illustrates the impact of increasing imbalance levels on the F1, precision, and recall measures. For the Compas and Adult Census datasets, which experience relatively more class and protected group imbalance, FOS shows greater resilience.

and installment credit features into second and third place. Feature weighting affects a classifier's decision boundary, which in turn, can impact fairness and accuracy. We can see in Figure 3 that FOS does change the placement of training examples. Figure 3 is a two dimensional depiction of the feature space based on t-SNE (Van der Maaten and Hinton, 2008). FOS generates additional synthetic examples of minority class, females (blue) and minority class, males (orange), as compared to the baseline image. We note that a sub-group of minority and majority females (blue and green) like closer to males (orange and red) and that there is also a small sub-group of males that lie closer to the female group in the FOS t-SNE image. Interestingly, both the baseline and FOS data show a more clear delineation between protected groups (females - blue and green vs. males - orange and red) than classes, which may imply underlying bias in the training data. For Adult Census, FOS changes the importance of the marital status, capital gain, age, sex and years of education features. See Figure 4. Based on its t-SNE, the arrangement of the Adult Census data appears more complex and less linearly separable. It should be noted that, based on its t-SNE, FOS appears to create more sub-groups of females (blue and green) that are embedded in the male cluster (6 clusters with FOS vs. only 2 in the baseline data). This additional sub-group creation may blur

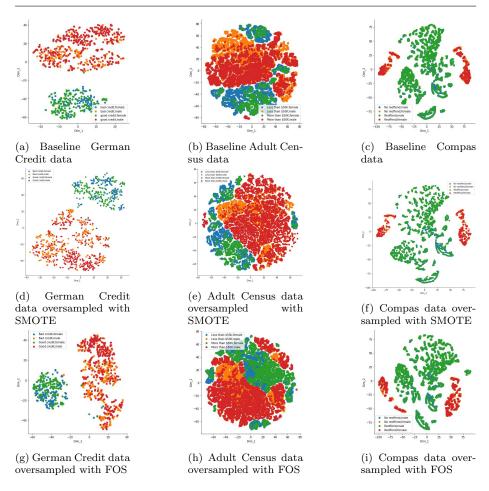


Fig. 3: The training datasets used in our experiments are split into four groups based on class and protected attributes and visualized with t-SNE. The under-represented groups are oversampled by FOS, resulting in more instances for the respective sensitive feature in the minority class. FOS reduces under-representation by protected groups and aids in blurring distinctions between sensitive features, thus reducing bias. See section 5.4 for discussion.

the separability between males and females, which causes a classifier to be less certain of the distinction between members of a protected group, thus reducing discrimination ($\mathbf{RQ4}$ answered).

7 Discussion and lessons learned

Discussion. Several papers have discussed the relationship between fairness and class imbalance in machine learning. Subramanian et al. note that there has traditionally been a disconnect between imbalanced learning and bias reduction; al-

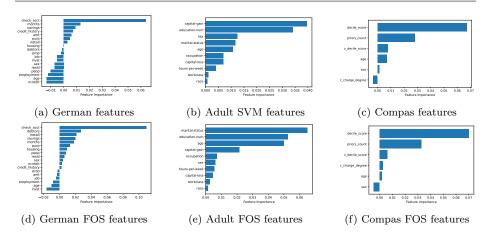


Fig. 4: Comparison of SVM feature importance with and without FOS. In the case of the German and Adult datasets, FOS induces a change in feature selection.

though they believe that the two areas share many similarities (Subramanian et al, 2021). They propose a margin-loss approach for tweet sentiment and occupation classification. Ferrari and Bacciu propose a cost-sensitive approach to deal with both imbalanced data and fairness (Ferrari and Bacciu, 2021). Iosifidis and Ntoutsi discuss imbalanced learning and fairness in the data streaming context (Iosifidis and Ntoutsi, 2020) and also consider using data augmentation techniques to increase the number of training instances for under-represented protected groups (Iosifidis and Ntoutsi, 2018).

Lessons learned. Our survey of existing approaches for algorithmic fairness and imbalanced data, as well as experiments conducted to evaluate the proposed FOS approach lead us to learn the following lessons:

- Two sides of the same coin. Fairness and class imbalance focus on targeting one specific type of bias present either in data or algorithms, addressing similar issues from different perspectives. However, these two domains are not disjoint and in fact should be considered as mutually complementary. Training data can be imbalanced with respect to classes and protected features. For example, the German Credit, Adult Census and Compas datasets (see Table 2) all exhibit imbalance with respect to classes and sensitive groups. Discrimination can result from the under-representation of protected groups in training data.
- Data imbalance techniques can help make standard algorithms more fair. Our experimental study shows that using a technique derived from the imbalanced domain (oversampling) with modifications ensuring its focus on fairness (see FOS description) can make standard classifiers fairness-aware much more efficiently than using several existing preprocessing techniques taken from either imbalanced or fairness domains. This shows the massive potential in bridging those two domains and creating approaches that leverage and merge state-of-the-art advancements in each of them.
- Data imbalance techniques can further boost fairness-aware algorithms. Our experiments additionally proved that by using techniques derived from imbalanced domain, we can significantly improve existing fairness-aware

- classifiers. This shows that multiple biases may exist in data at once and only by addressing the ways in which both data and algorithms may be skewed can we obtain truly fair machine learning.
- Robustness and fairness should be seen as a holistic task. Based on our previous observation, we postulate that robustness and fairness should be seen as a holistic task. When evaluating algorithms, we need to use practices from both domains. Our experiments show the importance of using fairness metrics combined with data-level difficulties, such as an increasing imbalance ratio. Additionally, understanding the impact of features on the classifier's decisions is crucial to evaluating if protected features truly have no impact on a classifier. This can also be seen as a first step towards explainable systems that shed light on both fairness and class imbalance.

8 Open challenges and future directions

Next, we formulate the open challenges and future directions that may be taken by researchers to further develop algorithms that are fair and robust by merging algorithmic fairness with class imbalance approaches.

- Unified nomenclature for fairness and imbalance. Based on our review of the fairness and imbalanced learning literature, we can see that although some of the nomenclature varies between these fields, the basic methodology is similar. For example, both imbalanced learning and fairness utilize three basic approaches: (1) pre-processing / data-level methods, (2) in-processing / algorithm level methods, and (3) post-processing / hybrid methods. Common metrics used in imbalanced learning and group fairness also are derived from a binary confusion matrix (Johnson and Khoshgoftaar, 2019). Therefore, there is a need for unified nomenclature accepted by researchers from both fields that will enable and foster collaboration between these two domains.
- Holistic and reproducible experimental framework. The algorithmic fairness domain could significantly benefit from incorporating similar rigorous experimental protocols that are nowadays becoming standard in imbalanced learning (Fernández et al, 2018). There is a need for agreed upon usage of benchmarks, performance metrics, baseline models, as well as verifying obtained results with statistical analysis (Stapor et al, 2021). By creating a widely accepted experimental testbed, the new methods that aim at countering various biases in data can be truly fairly analyzed and compared with state-of-the-art. Furthermore, a software library offering easy to use and an extensible set of the most effective methods for fairness and imbalance-aware methods should be developed.
- Exploration of different levels of interplay between fairness and imbalance domains. In this work, we showed how oversampling can be used not only to achieve fairness in classification, but also further boost fairness-aware classifiers. As the imbalanced learning literature is rich with diverse ideas on how to counter bias in algorithms(Fernández et al, 2018), we anticipate a great potential in applying other families of imbalanced approaches to algorithmic fairness. Cost-sensitive (Zhang et al, 2019a) and skew-insensitive (Wen and Wu, 2021) mechanisms stand out as the most obvious choices, but

we anticipate that other methods may prove beneficial, such as analyzing the usefulness of instance-level properties of minority classes and their impact on fairness (Skryjomski and Krawczyk, 2017), learning discriminative and fair low-dimensional representations (Korycki and Krawczyk, 2021b), or using one-class classification tailored for protected features (Bellinger et al, 2017).

- Robustness to adversarial attacks. The lack of fairness in machine learning may not only be caused by data or algorithms themselves. Decision-making systems may be subject to adversarial attacks caused by a malicious party that aims at damaging fairness to either discredit a given system or counter the progress towards more transparent and just society (Yeom and Fredrikson, 2020). We see developing robust mechanisms for fairness-aware algorithms as a crucial step towards fighting bias (Xu et al, 2021). This calls for understanding how fairness-aware algorithms can be poisoned by attacks, creating unified benchmarks with adversarial attacks targeting fairness, and developing adversarial training protocols to prepare robust learning systems to be deployed in the field.
- Fairness and robustness to imbalance in continual and streaming environments. All of the works discussed in this paper considered static classification problems. However, modern data sources constantly generate new information, leading to the need for learning from data streams and continual learning. The former focuses on adaptation to changes, known as concept drift (Korycki and Krawczyk, 2021a), while the latter focuses on retaining previously learned information and avoiding catastrophic forgetting (Parisi et al, 2019). Considering algorithmic fairness in such dynamic scenarios can lead to a plethora of exciting challenges, such as the evolving nature of bias in data, the emergence of new protected attributes over time, or selective forgetting and adaptation towards boosting fairness and reducing bias over time (Aguiar et al, 2022).
- XAI for fairness and robustness to imbalance. Explainable artificial intelligence allows end-users to understand the decision-making process of a classifier, as well as gain insights and better understanding of the analyzed data properties (Lin et al, 2021). This is crucial from the perspective of algorithmic fairness, as XAI should lead to transparency with respect to how a classifier was able to avoid bias (de Greeff et al, 2021). Furthermore, XAI may allow for discovering new, unknown biases in both data and the learning model itself (Fu et al, 2020). XAI has not been sufficiently explored in the context of imbalanced data, thus making it an important open challenge. Creating tools for the explainable understanding of various types of bias will allow us to further bridge these two domains by better comprehending the interplay between them.

9 Conclusion

A key facet in reducing algorithmic discrimination with respect to protected groups is the simultaneous reduction of class and protected group imbalance in training data. We show that reducing data imbalance facilitates improvements in both model accuracy and group fairness. This paper discussed the importance of bridging imbalanced learning and group fairness, by showing how key concepts in these

fields overlap, and it proposed a novel oversampling algorithm, Fair Oversampling, that addresses both skewed class distributions and protected attributes. Our method: (i) can be used as an efficient pre-processing algorithm with standard ML algorithms to jointly improve imbalance and fairness; and (ii) can supplement fairness-aware learning algorithms to improve their robustness. Additionally, we take a step toward bridging the gap between fairness and imbalanced learning with a new metric, Fair Utility, that combines balanced accuracy with group fairness measures.

References

- Agarwal A, Beygelzimer A, Dudík M, Langford J, Wallach H (2018) A reductions approach to fair classification. In: International Conference on Machine Learning, PMLR, pp 60–69
- Aguiar G, Krawczyk B, Cano A (2022) A survey on learning from imbalanced data streams: taxonomy, challenges, empirical study, and reproducible experimental framework. CoRR abs/2204.03719, DOI 10.48550/arXiv.2204.03719, URL https://doi.org/10.48550/arXiv.2204.03719, 2204.03719
- Angwin J (2016) Machine bias: There's software used across the country to predict future criminals and its biased against blacks. URL https://github.com/Trusted-AI/AIF360/blob/master/aif360/data/raw/compas/README.md
- Barocas S, Selbst AD (2016) Big data's disparate impact. Calif L Rev 104:671 Bej S, Davtyan N, Wolfien M, Nassar M, Wolkenhauer O (2021) Loras: an oversampling approach for imbalanced datasets. Mach Learn 110(2):279–301
- Bellamy RK, Dey K, Hind M, Hoffman SC, Houde S, Kannan K, Lohia P, Martino J, Mehta S, Mojsilovic A, et al (2018) Ai fairness 360: An extensible toolkit for detecting, understanding, and mitigating unwanted algorithmic bias. arXiv preprint arXiv:181001943
- Bellinger C, Sharma S, Zaïane OR, Japkowicz N (2017) Sampling a longer life: Binary versus one-class classification revisited. In: First International Workshop on Learning with Imbalanced Domains: Theory and Applications, LIDTA@PKDD/ECML 2017, 22 September 2017, Skopje, Macedonia, PMLR, Proceedings of Machine Learning Research, vol 74, pp 64–78
- Berk R, Heidari H, Jabbari S, Kearns M, Roth A (2021) Fairness in criminal justice risk assessments: The state of the art. Sociological Methods & Research 50(1):3–44
- Biswas S, Rajan H (2020) Do the machine learning models on a crowd sourced platform exhibit bias? an empirical study on model fairness. In: Proceedings of the 28th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering, pp 642–653
- Bonchi F, Hajian S, Mishra B, Ramazzotti D (2017) Exposing the probabilistic causal structure of discrimination. International Journal of Data Science and Analytics 3(1):1–21
- Bunkhumpornpat C, Sinapiromsaran K, Lursinsap C (2009) Safe-level-smote: Safe-level-synthetic minority over-sampling technique for handling the class imbalanced problem. In: Pacific-Asia conference on knowledge discovery and data mining, Springer, pp 475–482

- Calders T, Verwer S (2010) Three naive bayes approaches for discrimination-free classification. Data mining and knowledge discovery 21(2):277–292
- Caliskan A, Bryson JJ, Narayanan A (2017) Semantics derived automatically from language corpora contain human-like biases. Science 356(6334):183–186
- Calmon FP, Wei D, Vinzamuri B, Ramamurthy KN, Varshney KR (2017) Optimized pre-processing for discrimination prevention. In: Proceedings of the 31st International Conference on Neural Information Processing Systems, pp 3995–4004
- Cao K, Wei C, Gaidon A, Arechiga N, Ma T (2019) Learning imbalanced datasets with label-distribution-aware margin loss. arXiv preprint arXiv:190607413
- Chakraborty J, Majumder S, Menzies T (2021) Bias in machine learning software: Why? how? what to do? arXiv preprint arXiv:210512195
- Chawla NV, Bowyer KW, Hall LO, Kegelmeyer WP (2002) Smote: synthetic minority over-sampling technique. Journal of artificial intelligence research 16:321–357
- Chouldechova A (2017) Fair prediction with disparate impact: A study of bias in recidivism prediction instruments. Big data 5(2):153-163
- Chouldechova A, Benavides-Prado D, Fialko O, Vaithianathan R (2018) A case study of algorithm-assisted decision making in child maltreatment hotline screening decisions. In: Conference on Fairness, Accountability and Transparency, PMLR, pp 134–148
- Cieslak DA, Chawla NV, Striegel A (2006) Combating imbalance in network intrusion datasets. In: GrC, Citeseer, pp 732–737
- Commission E (2021) Proposal for a regulation laying down harmonised rules on artificial intelligence (artificial intelligence act). URL https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX%3A52021PC0206
- Corbett-Davies S, Goel S (2018) The measure and mismeasure of fairness: A critical review of fair machine learning. arXiv preprint arXiv:180800023
- Corbett-Davies S, Pierson E, Feller A, Goel S, Huq A (2017) Algorithmic decision making and the cost of fairness. In: Proceedings of the 23rd acm sigkdd international conference on knowledge discovery and data mining, pp 797–806
- Cui Y, Jia M, Lin TY, Song Y, Belongie S (2019) Class-balanced loss based on effective number of samples. In: Proceedings of the IEEE/CVF conference on computer vision and pattern recognition, pp 9268–9277
- Dablain D, Krawczyk B, Chawla NV (2022) DeepSMOTE: Fusing deep learning and smote for imbalanced data. IEEE Transactions on Neural Networks and Learning Systems pp 1–15, DOI 10.1109/TNNLS.2021.3136503
- Dastin J (2018)Amazon scraps secret ai recruiting tool https://www. that showed bias against women. URL reuters.com/article/us-amazon-com-jobs-automation-insight/
 - amazon-scraps-secret-ai-recruiting-tool-that-showed-bias-against-women-idUSKCN1MK08G
- Dwork C, Hardt M, Pitassi T, Reingold O, Zemel R (2012) Fairness through awareness. In: Proceedings of the 3rd innovations in theoretical computer science conference, pp 214–226
- Edwards H, Storkey A (2015) Censoring representations with an adversary. arXiv preprint arXiv:151105897
- Feldman M, Friedler SA, Moeller J, Scheidegger C, Venkatasubramanian S (2015) Certifying and removing disparate impact. In: proceedings of the 21th ACM SIGKDD international conference on knowledge discovery and data mining, pp

259-268

- Fernández A, García S, Galar M, Prati RC, Krawczyk B, Herrera F (2018) Learning from Imbalanced Data Sets. Springer, DOI 10.1007/978-3-319-98074-4, URL https://doi.org/10.1007/978-3-319-98074-4
- Ferrari E, Bacciu D (2021) Addressing fairness, bias and class imbalance in machine learning: the fbi-loss. arXiv preprint arXiv:210506345
- Flores AW, Bechtel K, Lowenkamp CT (2016) False positives, false negatives, and false analyses: A rejoinder to machine bias: There's software used across the country to predict future criminals. and it's biased against blacks. Fed Probation 80:38
- Fu Z, Xian Y, Gao R, Zhao J, Huang Q, Ge Y, Xu S, Geng S, Shah C, Zhang Y, de Melo G (2020) Fairness-aware explainable recommendation over knowledge graphs. In: Proceedings of the 43rd International ACM SIGIR conference on research and development in Information Retrieval, SIGIR 2020, Virtual Event, China, July 25-30, 2020, ACM, pp 69–78
- de Greeff J, de Boer MHT, Hillerström FHJ, Bomhof F, Jorritsma W, Neerincx MA (2021) The FATE system: Fair, transparent and explainable decision making. In: Proceedings of the AAAI 2021 Spring Symposium on Combining Machine Learning and Knowledge Engineering (AAAI-MAKE 2021), Stanford University, Palo Alto, California, USA, March 22-24, 2021, CEUR-WS.org, CEUR Workshop Proceedings, vol 2846
- Grgic-Hlaca N, Zafar MB, Gummadi KP, Weller A (2016) The case for process fairness in learning: Feature selection for fair decision making. In: NIPS symposium on machine learning and the law, vol 1, p 2
- Halevy A, Norvig P, Pereira F (2009) The unreasonable effectiveness of data. IEEE Intelligent Systems 24(2):8–12
- Han H, Wang WY, Mao BH (2005) Borderline-smote: a new over-sampling method in imbalanced data sets learning. In: International conference on intelligent computing, Springer, pp 878–887
- Hardt M, Price E, Srebro N (2016) Equality of opportunity in supervised learning. Advances in neural information processing systems 29:3315–3323
- Hashimoto T, Srivastava M, Namkoong H, Liang P (2018) Fairness without demographics in repeated loss minimization. In: International Conference on Machine Learning, PMLR, pp 1929–1938
- Hofmann DH (1994) Statlog (german credit data) data set. URL https://archive.ics.uci.edu/ml/datasets/Statlog+%28German+Credit+Data%29
- Iosifidis V, Ntoutsi E (2018) Dealing with bias via data augmentation in supervised learning scenarios. Jo Bates Paul D Clough Robert Jäschke 24
- Iosifidis V, Ntoutsi E (2020) Fabboo: Online fairness-aware learning under class imbalance. In: International Conference on Discovery Science, Springer, pp 159– 174
- Johnson JM, Khoshgoftaar TM (2019) Survey on deep learning with class imbalance. Journal of Big Data 6(1):1-54
- Kahn H, Marshall AW (1953) Methods of reducing sample size in monte carlo computations. Journal of the Operations Research Society of America 1(5):263–278
- Kamiran F, Žliobaite I, Calders T (2013) Quantifying explainable discrimination and removing illegal discrimination in automated decision making. Knowledge and information systems 35(3):613–644

- Khandani AE, Kim AJ, Lo AW (2010) Consumer credit-risk models via machine-learning algorithms. Journal of Banking & Finance 34(11):2767–2787
- Kleinberg J (2018) Inherent trade-offs in algorithmic fairness. In: Abstracts of the 2018 ACM International Conference on Measurement and Modeling of Computer Systems, pp 40–40
- Kohavi R, et al (1996) Scaling up the accuracy of naive-bayes classifiers: A decision-tree hybrid. In: Kdd, vol 96, pp 202–207
- Korycki L, Krawczyk B (2021a) Class-incremental experience replay for continual learning under concept drift. In: IEEE Conference on Computer Vision and Pattern Recognition Workshops, CVPR Workshops 2021, virtual, June 19-25, 2021, Computer Vision Foundation / IEEE, pp 3649–3658
- Korycki L, Krawczyk B (2021b) Low-dimensional representation learning from imbalanced data streams. In: Advances in Knowledge Discovery and Data Mining 25th Pacific-Asia Conference, PAKDD 2021, Virtual Event, May 11-14, 2021, Proceedings, Part I, Springer, Lecture Notes in Computer Science, vol 12712, pp 629-641
- Krawczyk B (2016) Learning from imbalanced data: open challenges and future directions. Progress in Artificial Intelligence 5(4):221–232
- Krawczyk B, Koziarski M, Wozniak M (2020) Radial-based oversampling for multiclass imbalanced data classification. IEEE Trans Neural Networks Learn Syst 31(8):2818–2831
- Kubat M, Holte RC, Matwin S (1998) Machine learning for the detection of oil spills in satellite radar images. Machine learning 30(2):195–215
- Lango M, Stefanowski J (2018) Multi-class and feature selection extensions of roughly balanced bagging for imbalanced data. J Intell Inf Syst 50(1):97–127
- Lin TY, Goyal P, Girshick R, He K, Dollár P (2017) Focal loss for dense object detection. In: Proceedings of the IEEE international conference on computer vision, pp 2980–2988
- Lin Y, Lee W, Celik ZB (2021) What do you see?: Evaluation of explainable artificial intelligence (XAI) interpretability through neural backdoors. In: Zhu F, Ooi BC, Miao C (eds) KDD '21: The 27th ACM SIGKDD Conference on Knowledge Discovery and Data Mining, Virtual Event, Singapore, August 14-18, 2021, ACM, pp 1027–1035
- Van der Maaten L, Hinton G (2008) Visualizing data using t-sne. Journal of machine learning research 9(11)
- Narayanan A (2018) Translation tutorial: 21 fairness definitions and their politics. In: Proc. Conf. Fairness Accountability Transp., New York, USA, vol 1170
- Parisi GI, Kemker R, Part JL, Kanan C, Wermter S (2019) Continual lifelong learning with neural networks: A review. Neural Networks 113:54–71
- Quy TL, Roy A, Iosifidis V, Ntoutsi E (2021) A survey on datasets for fairness-aware machine learning. CoRR abs/2110.00530, URL https://arxiv.org/abs/2110.00530, 2110.00530
- Raji ID, Gebru T, Mitchell M, Buolamwini J, Lee J, Denton E (2020) Saving face: Investigating the ethical concerns of facial recognition auditing. In: Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society, pp 145–151
- Rao RB, Krishnan S, Niculescu RS (2006) Data mining for improved cardiac care. Acm Sigkdd Explorations Newsletter 8(1):3–10
- Rawls J (2001) Justice as fairness: A restatement. Harvard University Press

- Ridnik T, Ben-Baruch E, Zamir N, Noy A, Friedman I, Protter M, Zelnik-Manor L (2021) Asymmetric loss for multi-label classification. In: Proceedings of the IEEE/CVF International Conference on Computer Vision, pp 82–91
- Romei A, Ruggieri S (2014) A multidisciplinary survey on discrimination analysis. The Knowledge Engineering Review 29(5):582–638
- Schumann C, Foster J, Mattei N, Dickerson J (2020) We need fairness and explainability in algorithmic hiring. In: International Conference on Autonomous Agents and Multi-Agent Systems (AAMAS)
- Seiffert C, Khoshgoftaar TM, Van Hulse J, Napolitano A (2007) Mining data with rare events: a case study. In: 19th IEEE International Conference on Tools with Artificial Intelligence (ICTAI 2007), IEEE, vol 2, pp 132–139
- Sharma S, Bellinger C, Krawczyk B, Zaïane OR, Japkowicz N (2018) Synthetic oversampling with the majority class: A new perspective on handling extreme imbalance. In: IEEE International Conference on Data Mining, ICDM 2018, Singapore, November 17-20, 2018, IEEE Computer Society, pp 447–456
- Shroff R (2017) Predictive analytics for city agencies: Lessons from children's services. Big data 5(3):189-196
- Skryjomski P, Krawczyk B (2017) Influence of minority class instance types on SMOTE imbalanced data oversampling. In: First International Workshop on Learning with Imbalanced Domains: Theory and Applications, LIDTA@PKDD/ECML 2017, 22 September 2017, Skopje, Macedonia, PMLR, Proceedings of Machine Learning Research, vol 74, pp 7–21
- Sleeman WC, Krawczyk B (2021) Multi-class imbalanced big data classification on spark. Knowl Based Syst 212:106598
- Speicher T, Heidari H, Grgic-Hlaca N, Gummadi KP, Singla A, Weller A, Zafar MB (2018) A unified approach to quantifying algorithmic unfairness: Measuring individual &group unfairness via inequality indices. In: Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining, pp 2239–2248
- Stapor K, Ksieniewicz P, García S, Wozniak M (2021) How to design the fair experimental classifier evaluation. Appl Soft Comput 104:107219
- Subramanian S, Rahimi A, Baldwin T, Cohn T, Frermann L (2021) Fairness-aware class imbalanced learning. arXiv preprint arXiv:210910444
- Tao X, Li Q, Guo W, Ren C, Li C, Liu R, Zou J (2019) Self-adaptive cost weights-based support vector machine cost-sensitive ensemble for imbalanced data classification. Inf Sci 487:31–56
- Wei W, Li J, Cao L, Ou Y, Chen J (2013) Effective detection of sophisticated online banking fraud on extremely imbalanced data. World Wide Web 16(4):449–475
- Wen G, Wu K (2021) Building decision tree for imbalanced classification via deep reinforcement learning. In: Asian Conference on Machine Learning, ACML 2021, 17-19 November 2021, Virtual Event, PMLR, Proceedings of Machine Learning Research, vol 157, pp 1645–1659
- Woźniak M, Grana M, Corchado E (2014) A survey of multiple classifier systems as hybrid systems. Information Fusion 16:3–17
- Xu H, Liu X, Li Y, Jain AK, Tang J (2021) To be robust or to be fair: Towards fairness in adversarial training. In: Proceedings of the 38th International Conference on Machine Learning, ICML 2021, 18-24 July 2021, Virtual Event, PMLR, Proceedings of Machine Learning Research, vol 139, pp 11492–11501

- Yeom S, Fredrikson M (2020) Individual fairness revisited: Transferring techniques from adversarial robustness. In: Proceedings of the Twenty-Ninth International Joint Conference on Artificial Intelligence, IJCAI 2020, ijcai.org, pp 437–443
- Yurochkin M, Bower A, Sun Y (2019) Training individually fair ml models with sensitive subspace robustness. arXiv preprint arXiv:190700020
- Zafar MB, Valera I, Rogriguez MG, Gummadi KP (2017) Fairness constraints: Mechanisms for fair classification. In: Artificial Intelligence and Statistics, PMLR, pp 962–970
- Zemel R, Wu Y, Swersky K, Pitassi T, Dwork C (2013) Learning fair representations. In: International conference on machine learning, PMLR, pp 325–333
- Zhang BH, Lemoine B, Mitchell M (2018) Mitigating unwanted biases with adversarial learning. In: Proceedings of the 2018 AAAI/ACM Conference on AI, Ethics, and Society, pp 335-340
- Zhang C, Tan KC, Li H, Hong GS (2019a) A cost-sensitive deep belief network for imbalanced classification. IEEE Trans Neural Networks Learn Syst 30(1):109–122
- Zhang W, Ramezani R, Naeim A (2019b) Wotboost: Weighted oversampling technique in boosting for imbalanced learning. In: 2019 IEEE International Conference on Big Data (IEEE BigData), Los Angeles, CA, USA, December 9-12, 2019, IEEE, pp 2523–2531
- Žliobaite I (2017) Measuring discrimination in algorithmic decision making. Data Mining and Knowledge Discovery 31(4):1060–1089
- Zyblewski P, Sabourin R, Wozniak M (2021) Preprocessed dynamic classifier ensemble selection for highly imbalanced drifted data streams. Inf Fusion 66:138–154